

# Curso Data Engineer: Creando un pipeline de datos

MÓDULO E - Clase 7

# Agenda



- Transformaciones en cloud (Spark)
- Orquestación
- Ejercitación




# Transformaciones en cloud


# Data Proc




Row data



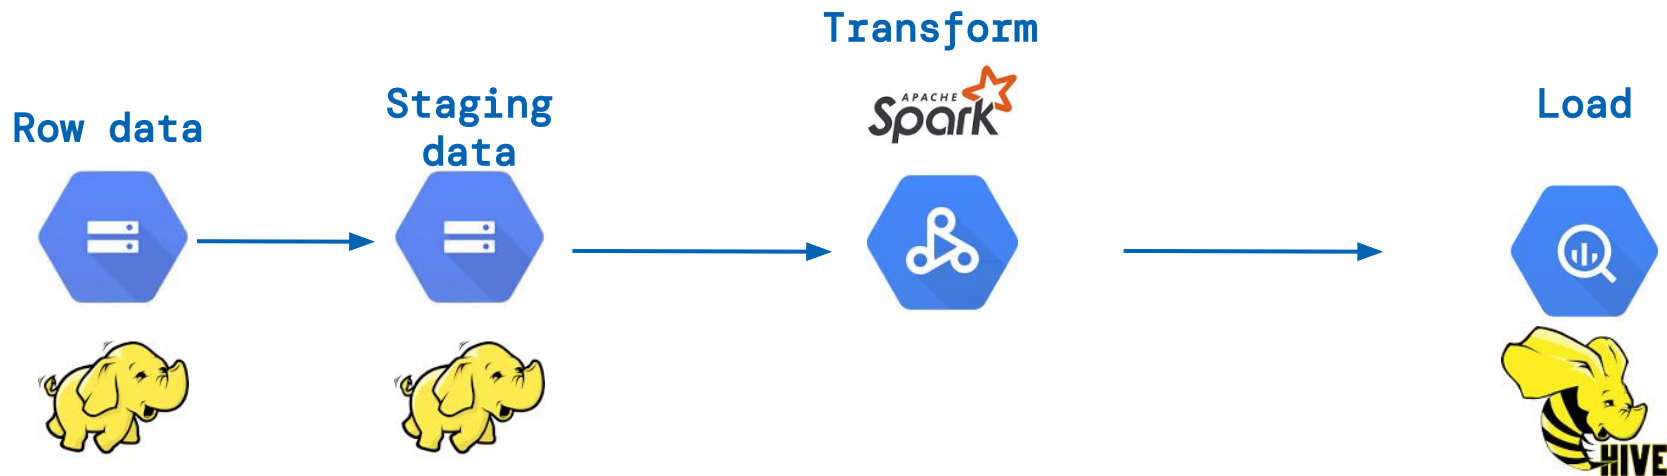
# Data Proc



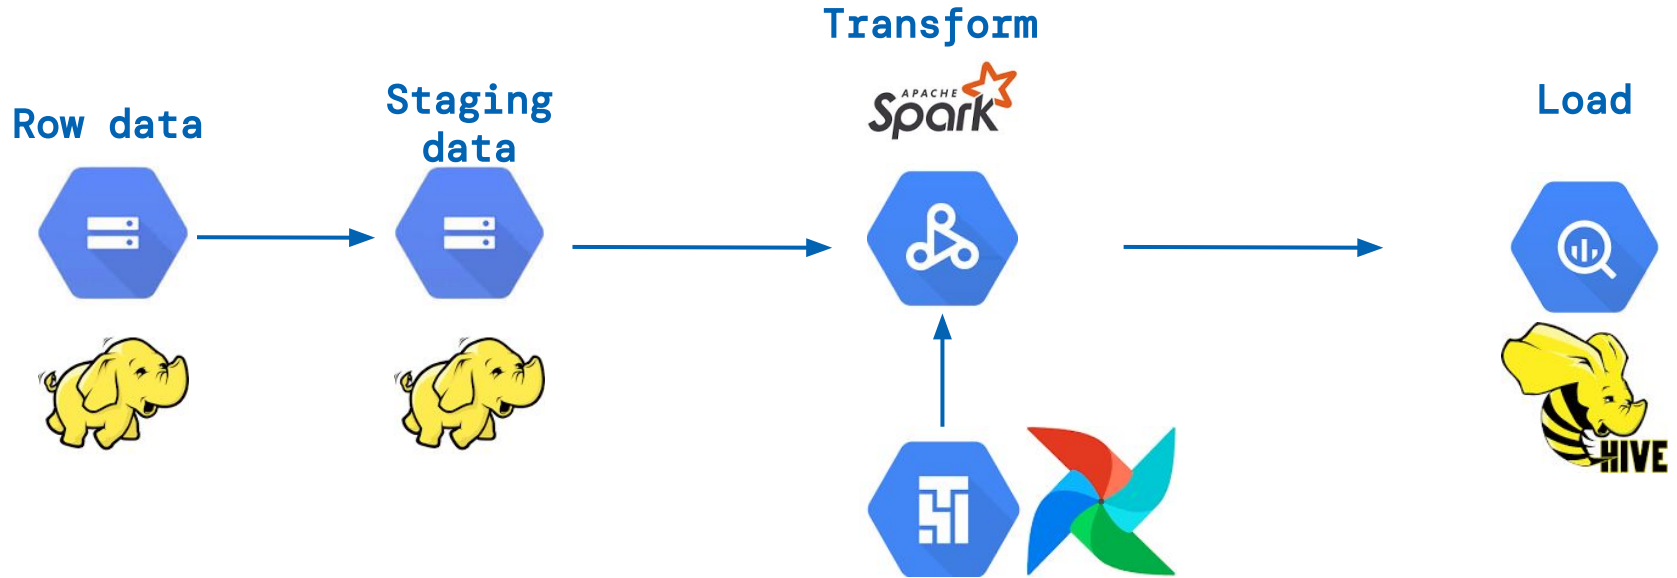
# Data Proc



# Data Proc

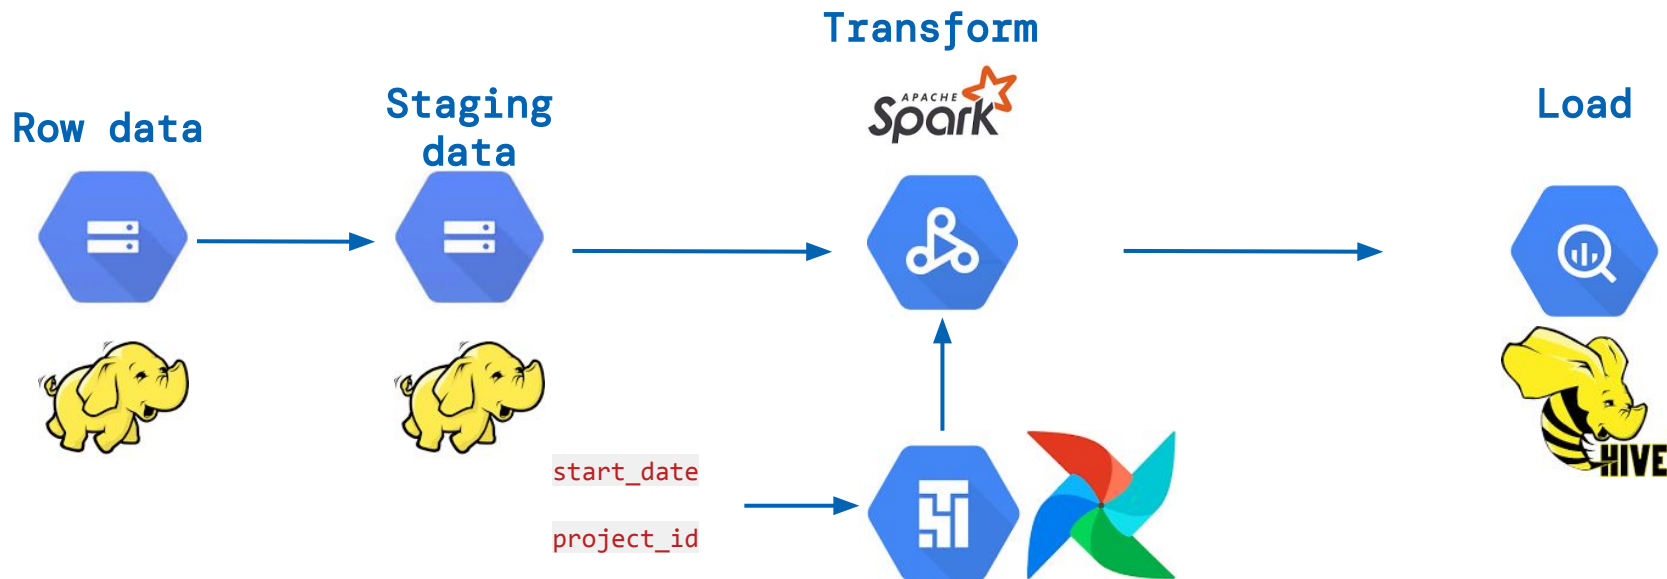


# Data Proc





# Data Proc



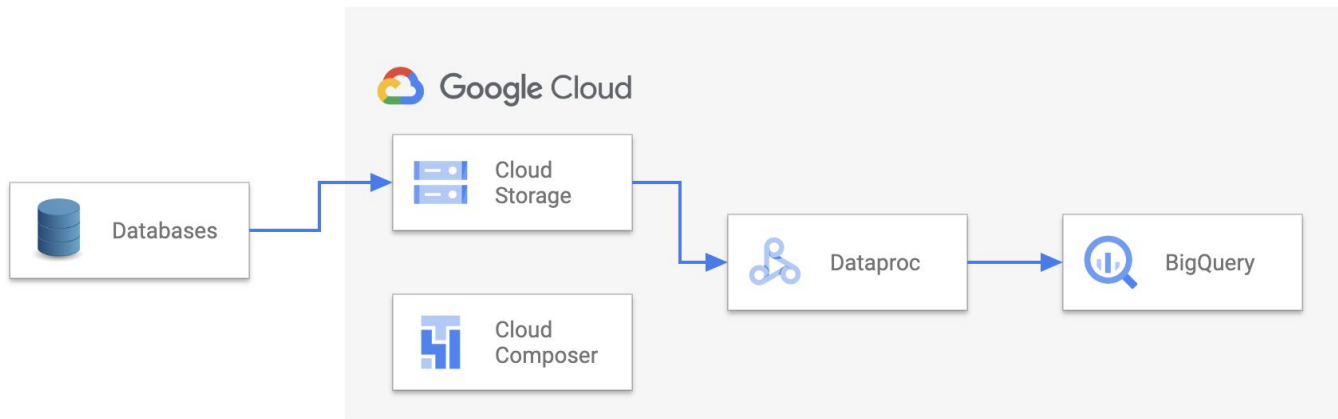


# Orquestación

# Batch process




Batch



# Streaming process

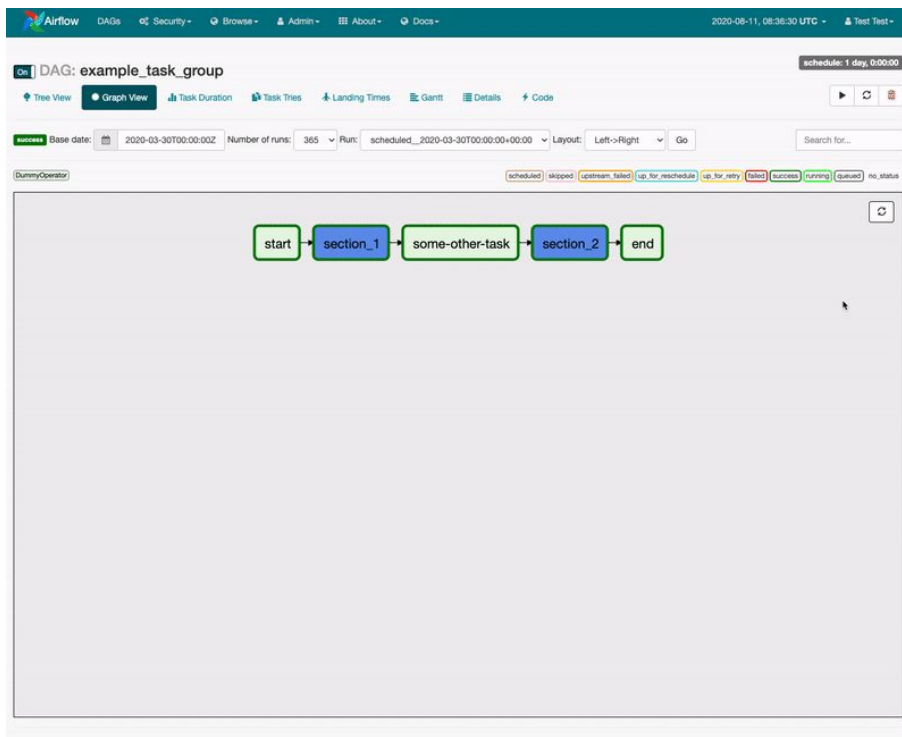


Streaming



# Apache Airflow

REAL \* IA PARA  
UN MUNDO



# Apache Airflow

## Importing Modules

An Airflow pipeline is just a Python script that happens to define an Airflow DAG object. Let's start by importing the libraries we will need.

airflow/example\_dags/tutorial.py

[\[source\]](#)

```
from datetime import datetime, timedelta
from textwrap import dedent

# The DAG object; we'll need this to instantiate a DAG
from airflow import DAG

# Operators; we need this to operate!
from airflow.operators.bash import BashOperator
```

# Apache Airflow

## Default Arguments

We're about to create a DAG and some tasks, and we have the choice to explicitly pass a set of arguments to each task's constructor (which would become redundant), or (better!) we can define a dictionary of default parameters that we can use when creating tasks.

REAL \* IA PARA  
UN MUNDO

airflow/example\_dags/tutorial.py

[\[source\]](#)

```
# These args will get passed on to each operator
# You can override them on a per-task basis during operator initialization
default_args={
    "depends_on_past": False,
    "email": ["airflow@example.com"],
    "email_on_failure": False,
    "email_on_retry": False,
    "retries": 1,
    "retry_delay": timedelta(minutes=5),
    # 'queue': 'bash_queue',
    # 'pool': 'backfill',
    # 'priority_weight': 10,
    # 'end_date': datetime(2016, 1, 1),
    # 'wait_for_downstream': False,
    # 'sla': timedelta(hours=2),
    # 'execution_timeout': timedelta(seconds=300),
    # 'on_failure_callback': some_function,
    # 'on_success_callback': some_other_function,
    # 'on_retry_callback': another_function,
    # 'sla_miss_callback': yet_another_function,
    # 'trigger_rule': 'all_success'
},
```

# Apache Airflow

## Instantiate a DAG

We'll need a DAG object to nest our tasks into. Here we pass a string that defines the `dag_id`, which serves as a unique identifier for your DAG. We also pass the default argument dictionary that we just defined and define a `schedule` of 1 day for the DAG.

REAL \* IA PARA UN MUNDO

airflow/example\_dags/tutorial.py

[\[source\]](#)

```
with DAG(
    "tutorial",
    # These args will get passed on to each operator
    # You can override them on a per-task basis during operator initialization
    default_args={
        "depends_on_past": False,
        "email": ["airflow@example.com"],
        "email_on_failure": False,
        "email_on_retry": False,
        "retries": 1,
        "retry_delay": timedelta(minutes=5),
        # 'queue': 'bash_queue',
        # 'pool': 'backfill',
        # 'priority_weight': 10,
        # 'end_date': datetime(2016, 1, 1),
        # 'wait_for_downstream': False,
        # 'sla': timedelta(hours=2),
        # 'execution_timeout': timedelta(seconds=300),
        # 'on_failure_callback': some_function,
        # 'on_success_callback': some_other_function,
        # 'on_retry_callback': another_function,
        # 'sla_miss_callback': yet_another_function,
        # 'trigger_rule': 'all_success'
    },
    description="A simple tutorial DAG",
    schedule=timedelta(days=1),
    start_date=datetime(2021, 1, 1),
    catchup=False,
    tags=["example"],
) as dag:
```



# Apache Airflow

## Operators

An Operator is conceptually a template for a predefined [Task](#), that you can just define declaratively inside your DAG:

```
with DAG("my-dag") as dag:
    ping = SimpleHttpOperator(endpoint="http://example.com/update/")
    email = EmailOperator(to="admin@example.com", subject="Update complete")

    ping >> email
```

Airflow has a very extensive set of operators available, with some built-in to the core or pre-installed providers. Some popular operators from core include:

- **BashOperator** - executes a bash command
- **PythonOperator** - calls an arbitrary Python function
- **EmailOperator** - sends an email

# Apache Airflow



If the operator you need isn't installed with Airflow by default, you can probably find it as part of our huge set of community [provider packages](#). Some popular operators from here include:

- `SimpleHttpOperator`
- `MySqlOperator`
- `PostgresOperator`
- `MsSqlOperator`
- `OracleOperator`
- `JdbcOperator`
- `DockerOperator`
- `HiveOperator`
- `S3FileTransformOperator`
- `PrestoToMySqlOperator`
- `SlackAPIOperator`

# Apache Airflow



## Tasks

To use an operator in a DAG, you have to instantiate it as a task. Tasks determine how to execute your operator's work within the context of a DAG.

In the following example, we instantiate the BashOperator as two separate tasks in order to run two separate bash scripts. The first argument for each instantiation, `task_id`, acts as a unique identifier for the task.

airflow/example\_dags/tutorial.py

[\[source\]](#)

```
t1 = BashOperator(
    task_id="print_date",
    bash_command="date",
)

t2 = BashOperator(
    task_id="sleep",
    depends_on_past=False,
    bash_command="sleep 5",
    retries=3,
)
```



# Ejercicio

# Apache Airflow

Username: airflow  
Password: airflow



The screenshot shows the Apache Airflow web interface. At the top, there is a teal header bar with the Airflow logo on the left, the date and time '2020-05-03, 15:19:11 UTC' in the center, and a 'Login' link on the right. Below the header, there is a 'Sign In' modal box. Inside the modal, it says 'Enter your login and password below:'. There are two input fields: 'Username:' and 'Password:'. The 'Username:' field has a small person icon to its left, and the 'Password:' field has a small eye icon to its left. Below the input fields is a teal 'Sign In' button.

<http://<direccion ip>:8010>



# ejercicio

Crear un pipeline a través de un DAG, creando las tareas:

- Ingest
- Transform
- Load

# Load

[https://edvaibucket.blob.core.windows.net/data-engineer-edvai/yellow\\_tripdata\\_2021-01.csv?sp=r&st\=2023-11-06T12:52:39Z\&se\=2025-11-06T20:52:39Z\&sv\=2022-11-02\&sr\=c\&sig\=J4Ddi2c7Ep230hQLPisbYaerIH472iigPwc1%2FkG80EM%3D](https://edvaibucket.blob.core.windows.net/data-engineer-edvai/yellow_tripdata_2021-01.csv?sp=r&st\=2023-11-06T12:52:39Z\&se\=2025-11-06T20:52:39Z\&sv\=2022-11-02\&sr\=c\&sig\=J4Ddi2c7Ep230hQLPisbYaerIH472iigPwc1%2FkG80EM%3D)



Extract



Transform



Load



Orchestration

# Workaround possible issue

```
root@0485e86b7577:/# ping www.google.com
ping: www.google.com: Temporary failure in name resolution
```

```
root@0485e86b7577:/# wget -P /home/hadoop/landing/ https://data-engineer-edvai.s3.amazonaws.com/yellow_tripdata_2021-01.csv
--2023-04-17 18:52:31-- https://data-engineer-edvai.s3.amazonaws.com/yellow_tripdata_2021-01.csv
Resolving data-engineer-edvai.s3.amazonaws.com (data-engineer-edvai.s3.amazonaws.com)... failed: Temporary failure in name resolution.
wget: unable to resolve host address 'data-engineer-edvai.s3.amazonaws.com'
```

Dentro de consola\_Hadoop, modificar el archivo /etc/resolv.conf, comentado todos los DNS y dejando solamente el de Google, de la siguiente manera:

```
nameserver 8.8.8.8
options timeout:2 attempts:5
```